When the Simulator Is Wrong, Do Better Derivatives Improve Decisions?

Melvin Loh
Nanyang Technological University
Turing AI Scholars Programme
College of Computing and Data Science
mloh018@e.ntu.edu.sg

Prathamesh Dinesh Joshi
Vizuara AI Labs
prathamesh@vizuara.com

Abstract

We use option pricing and hedging to study whether better surrogate derivatives improve decisions when the assumed simulator is wrong. A neural Heston surrogate predicts option values. A fixed hedging rule uses their sensitivity to asset price to choose an asset holding intended to offset changes in the option's value. At two Heston coefficient regimes excluded from training and validation, adding curvature labels to price-sensitivity supervision reduces root-mean-square error (RMSE) in the unsupervised variance sensitivity by 40 and 26%. Supervising all three targets reduces curvature and variance-sensitivity RMSE by 82–93% relative to two baselines without derivative labels. The incremental gains depend on the target and are assessed against measured solver disagreement. When simulated dynamics depart from the Heston model class, a value-only model captures at least 96% of each of the fully supervised model's eight improvements over the residual-only baseline at zero transaction cost. At the registered 1% cost, the fully supervised model is worse than that baseline in all ten departure settings. The preregistered confirmatory hedging endpoint, evaluated under a coefficient shift, also fails. These findings support cross-target transfer within the tested nested design, without establishing a consistent hedging advantage specific to derivative labels or a general ordering of label choices.

1 Introduction

Scientific surrogates approximate models that can themselves be wrong. We use option pricing and hedging to ask whether reducing derivative error improves decisions under model misspecification. The network predicts a Heston option value; a fixed hedging rule uses its asset-price sensitivity. Numerical references measure value and derivative errors; simulations outside Heston test whether these errors matter for hedging. A close value fit does not ensure accurate derivatives: at a held-out coefficient regime, two models without derivative labels have value errors of 1.6% and 3.3%, but curvature errors of 41% and 27% (Appendix D). A physics residual also leaves substantial derivative error, consistent with documented limitations (Krishnapriyan et al., 2021; Wang et al., 2021).

Derivative supervision is established in Sobolev training (Czarnecki et al., 2017) and differential machine learning (Huge and Savine, 2020). Molecular models fit forces (Ercolessi and Adams, 1994; Wang et al., 2019) or differentiate learned energies (Behler and Parrinello, 2007). Sobolev-PINNs augment residual training (Son et al., 2023). Our contribution is a controlled evaluation of which targets improve and whether this improves hedging. Model misspecification also affects simulation-based inference (Cranmer et al., 2020; Cannon et al., 2022). Sequential prediction also highlights the importance of evaluation distributions (Ross et al., 2011).

40th Conference on Neural Information Processing Systems (NeurIPS 2026). Workshop: Sim2Science: ML with Imperfect Scientific Models.

Does fidelity translate under misspecification? We distinguish held-out coefficients within the assumed model from a generating process outside its class, which no parameter setting recovers. In the financial setting, this means asking whether more accurate Heston derivatives improve hedging when the simulated price process departs from Heston dynamics.

Where should the labels go? Derivative labels require additional reference computations. We add targets sequentially to measure their conditional contribution. Noguer i Alonso (2026) identifies the importance of second derivatives in discrete-time hedging but reports no hedging backtest. Malandain et al. (2025) reports inaccurate derivatives despite a small residual and proposes derivative regularisation. The closest trained configuration (Sakuma, 2026) supervises first and second derivatives together, but evaluates in distribution and reports similar pricing and Greek accuracy among its three supervised arms.

We compare a nested sequence of label sets, pre-registered controls, and numerical references with measured uncertainty. The registered tests and thresholds were fixed before training. Numerical cross-checks let us measure errors in the value and the evaluated derivatives, and assess whether differences exceed reference uncertainty.

What we find. Full supervision reduces held-out curvature and variance-derivative RMSE by 82–93% relative to the value-only and residual-only baselines. Adding curvature labels also improves the unlabelled variance derivative by 40 and 26%, but adjacent-rung gains depend on the target. Under model-class departure, value-only training largely reproduces the zero-cost hedging gains; at 1% cost the fully supervised model is worse than the residual-only baseline in all ten settings. The registered confirmatory test also fails (Sections 3.1 and 3.6).

What this is not. We do not correct the simulator's dynamics. We reduce the surrogate's error against an assumed pricing model and ask whether this improves a fixed hedging rule when the model is structurally wrong. The experiments test this distinction in one PDE setting; they do not establish transfer to other simulators or real-market performance.

2 Setup

Model and notation. The Heston model (Heston, 1993) describes an asset price S and its instantaneous variance v, both evolving randomly. The function $u(S, v, \tau; K)$ is the model value of a European call option: a contract giving its holder the right to buy the asset at a fixed price K (the strike) at expiry. With time to expiry τ, its payoff at expiry is $\max(S - K, 0)$. For fixed coefficients, the pricing PDE is

$$\begin{aligned} \partial_\tau u = \tfrac{1}{2} v S^2 \partial_{SS} u + \rho \xi v S \partial_{Sv} u + \tfrac{1}{2} \xi^2 v \partial_{vv} u \\ + (r - q) S \partial_S u + \kappa(\theta - v) \partial_v u - r u, \end{aligned} \tag{1}$$

with $u(S, v, 0; K) = \max(S - K, 0)$. Here κ is the variance mean-reversion rate, θ its long-run level, ξ the variance diffusion scale, and ρ the correlation between the price and variance shocks; r and q are the interest rate and dividend yield. The variance input is denoted v_0 in the data and configuration. The network approximates u across coefficients, as in parametric Heston PINNs (Hainaut and Casas, 2024); the squared PDE residual measures its failure to satisfy Eq. (1) and supplies the physics-informed loss (Raissi et al., 2019).

The three targets are *delta*, the price sensitivity $\partial_S u$; *gamma*, the curvature $\partial_{SS} u = \partial^2 u / \partial S^2$; and the variance sensitivity $\partial_v u$. The last is a derivative with respect to variance, rather than volatility. Delta describes how the option value changes with asset price; gamma describes how that sensitivity itself changes. All are input derivatives obtained by automatic differentiation, distinct from loss gradients with respect to network weights. The registered hedging rule uses delta to set its asset holding; gamma and variance sensitivity are supervision targets, not additional inputs to that rule.

Reference cross-checks. We cross-check option values using up to four solvers: Heston Fourier inversion (Heston, 1993; Albrecher et al., 2007), cosine expansion (Fang and Oosterlee, 2008), an alternating-direction finite-difference scheme (Craig and Sneyd, 1988), and Monte Carlo (Lord et al., 2010). The first two are separately implemented but share a characteristic-function formula, so they cannot exclude a common formula error. The predeclared mask combines agreement checks across values and derivatives, with a relative tolerance of 10^{-3}, a scale floor near zero, and an allowance for Monte-Carlo noise (Appendix B). It removes 17.1% of the interpolated grid and 9.2% of the

extrapolated grid, concentrated at short expiries and just below the strike. For each metric, all models are evaluated on the same retained points.

What one envelope is. Value agreement does not certify derivative accuracy. For each derivative and evaluation regime, an *uncertainty envelope* combines the largest pairwise solver disagreement with Monte-Carlo standard error (Table 4). A difference is called *resolved* when it exceeds that envelope; this is a resolution heuristic, not a statistical test or error bound. Unresolved differences need not be zero. Labels and reference share a solver family: holding out parameters prevents memorisation of those regimes but cannot exclude shared bias. Appendix B quantifies these limits.

Departure. An *anchor* is a fixed coefficient setting excluded, together with a neighbourhood, from both training and validation. Derivative accuracy is evaluated at two Heston anchors: one inside the sampled coefficient ranges and one with correlation outside them. These test interpolation into a held-out region and extrapolation within the model class. For departures outside that class, we evaluate hedging under Bates dynamics, which add price jumps (Bates, 1996), and Merton dynamics, which combine jumps with constant volatility (Merton, 1976). In these tests every model receives the same assumed Heston coefficients; no derivative-accuracy claim is made for the departed dynamics (Appendix C).

Arms. An *arm* is one training configuration. Both baselines use value labels: the value-only model has no PDE residual, whereas the residual-only baseline adds that residual but no derivative labels. A nested *ladder* adds labels to the residual baseline: rung 1 adds delta, rung 2 also adds gamma, and rung 3 also adds variance sensitivity. Each label loss is a scale-normalised squared error (Huge and Savine, 2020). These arms share architecture, optimiser, data and budget. The residual, gamma and variance-label weights were selected on validation; the delta-label weight was fixed. Adaptive loss balancing was not tested. The extra-value-data control uses a different residual weight, disclosed in Section 3.5.

The hedging rule holds $\widehat{\partial_S u}$ units of the asset to offset changes in the option's value. It updates that holding over 43 steps on $10{,}000$ shared simulated trajectories (Andersen, 2008). Hedging loss J is expected shortfall, the mean of the worst 5% of losses, so it measures adverse outcomes rather than average performance (Rockafellar and Uryasev, 2000); lower is better. Tests include zero and nonzero transaction costs. The pre-registered thresholds require a 10% confirmatory improvement, a 15% derivative-accuracy improvement, and specify a 2% negligibility floor (Table 3).

3 Results

We report the registered intervals as specified before training. Later comparisons are exploratory, without adjustment for multiple comparisons. Path-bootstrap intervals measure variation across simulated paths for the fitted models, not variation across new training runs.

Table 1: Held-out derivative RMSE on the full parameter grid, interpolated / extrapolated coefficient regime, five seeds; second-derivative entries $\times 10^{-3}$. Lower is better. Bold marks a lowest mean separated from the runner-up by more than the reference uncertainty; an unbolded regime–derivative cell is one whose top two arms this benchmark cannot separate. Resolution is the registered reference-noise heuristic, not a test. Lower block: those contrasts divided by the reference envelope for that derivative; italics mark a gap under one envelope, which this benchmark cannot resolve (Appendix B).

Arm	$\partial u/\partial S$	$\partial^2 u/\partial S^2$	$\partial u/\partial v$
residual-only baseline	$0.0537/0.0424$	$7.82/6.04$	$9.93/7.44$
value-only feedforward	$0.0881/0.0815$	$8.29/7.17$	$5.74/2.89$
rung 1, $+\partial u/\partial S$	$0.0106/0.0079$	$2.56/1.58$	$2.66/1.73$
rung 2, $+\partial^2 u/\partial S^2$	$0.0033/0.0034$	$1.02/0.64$	$1.59/1.28$
rung 3, $+\partial u/\partial v$	$\mathbf{0.0027}/0.0038$	$0.83/0.82$	$\mathbf{0.73/0.52}$
same contrasts, in reference envelopes			
baseline → rung 3	$184.7/97.3$	$9.1/2.6$	$85.0/54.5$
rung 1 → rung 2	$26.2/11.4$	$2.0/\mathit{0.5}$	$9.9/3.5$
rung 2 → rung 3	$2.2/\mathit{-0.97}$	$\mathit{0.2/-0.1}$	$8.0/6.0$

3.1 Derivative accuracy and hedging under model departure

Table 1 evaluates derivative accuracy at held-out coefficients within the Heston class. Relative to the two baselines without derivative labels, full supervision reduces gamma and variance-sensitivity RMSE by 82–93% and delta RMSE by 91–97%. These baseline-to-rung-3 improvements exceed the reference uncertainty at both regimes (Appendix B). The derivative-accuracy findings below concern these two regimes; departures outside the model class are evaluated through hedging loss.

At zero transaction cost, rung 3 has lower mean loss than the residual-only baseline in seven of nine Bates jump settings and in the Merton setting. The value-only model captures at least 96% of each of these eight gains and outperforms rung 3 in four. At the registered 1% cost, rung 3 has higher loss than the residual-only baseline in all ten settings. Thus the observed gains over that baseline do not establish a consistent hedging advantage specific to derivative labels.

Extent of misspecification. The Bates sweep varies jump arrival rate $\lambda_j \leq 0.5$ and jump-size standard deviation $\sigma_j \leq 0.15$ (Appendix C). These parameters order the tested departures, but we did not measure distance from the Heston model class. At zero cost the relative gain over the residual-only baseline decreases along each coordinate, from $+27.7\%$ at the in-class setting to -4.2% at the largest tested jumps, changing sign along the $\lambda_j = 0.5$ row. All six one-dimensional chains decrease in every one of five seeds, for both absolute and relative gains ($30/30$ each). This monotonicity analysis is exploratory and differs from the preregistered expectation of a non-monotone pattern. The registered departure-sweep verdict uses the 1% cost tier, where no favourable regime is found.

The registered derivative-accuracy endpoint also fails, despite gamma and variance-sensitivity reductions of 86.4–93.0% exceeding its 15% threshold. It additionally requires value accuracy to remain within a two-sided parity band; improving value error by 82–91% violates that band. We retain the registered failure and disclose this limitation of the criterion.

3.2 Adding curvature labels improves curvature accuracy

Adding gamma labels to delta supervision reduces mean gamma RMSE by 60% and 59%, while value error changes by 4.4% at both regimes. The relative curvature improvement is therefore much larger than the change in value fit in this comparison. The gamma improvement is 2.0 reference envelopes at the interpolated regime and 0.5 at the extrapolated regime: it is resolved only at the former. The larger baseline-to-rung-2 improvement is resolved at both. This conditional comparison neither establishes a general preference for gamma labels nor changes the registered parity failure.

3.3 The gain is not confined to the supervised derivative

Rung 2 adds only gamma labels to rung 1, yet RMSE in the unsupervised variance sensitivity $\partial_v u$ falls from $2.66/1.73$ to $1.59/1.28$, a 40% and 26% reduction. Delta error also falls two- to three-fold. The variance-sensitivity gaps are 9.9 and 3.5 reference envelopes, so cross-target transfer is resolved at both regimes. One possible explanation is coupling through the PDE residual, which contains both $\partial_{SS} u$ and $\partial_v u$. That mechanism remains untested; varying the residual and label terms independently would help evaluate it.

What the nested design establishes. This comparison measures the effect of adding gamma labels when delta labels and the shared residual are already present. It does not separate their interactions or establish that gamma is the best next target. Missing configurations include delta plus variance labels and gamma labels alone. These are targeted follow-ups; a full factorial design would require all combinations of the three label types. The transfer result therefore applies to this sequence of added labels.

3.4 The third target yields a resolved gain primarily on itself

Adding variance-sensitivity labels in rung 3 reduces that target's RMSE from $1.59/1.28$ to $0.73/0.52$, differences of 8.0 and 6.0 reference envelopes. Delta also improves at the interpolated regime by 2.2 envelopes, while the gamma improvement is unresolved. At extrapolation, mean delta RMSE increases from 0.003399 to 0.003784, and gamma RMSE in units of 10^{-3} from 0.641 to 0.823, but the differences are only 0.97 and 0.09 envelopes. We therefore cannot conclude that either derivative degrades. All three supervised arms have gamma errors below its reference uncertainty

there (Appendix B). Only the newly supervised target improves beyond the reference uncertainty at both regimes.

The registered label-noise experiment also fails to demonstrate a monotone dose-response. In the confirmatory hedging comparison, adding gamma labels to delta supervision increases loss by a statistically nonzero amount below the 2% negligibility floor. These are distinct downstream tests, reported with their original verdicts in Table 3; they do not identify the mechanism behind the derivative-accuracy results.

3.5 Controls for label information

The pre-registered controls use a curvature penalty without labels, shuffled labels, or additional value data. Against the baseline's held-out gamma relative RMSE of 0.410, the penalty gives 0.982 and shuffled labels 0.408; neither reproduces the supervised improvement. The penalty weight was not tuned. The price-only sweep uses $4{,}096$–$20{,}480$ rows, fewer than the main arms' $23{,}623$, and a residual weight 100 times larger. Its validation gamma relative RMSE changes from 0.647 to 0.596; doubling width gives 0.709 ± 0.061. Because these training conditions differ, the sweep cannot establish whether extra value labels can substitute for derivative labels (Appendix D).

3.6 Hedging performance and the registered endpoint

The confirmatory hedging test fails: the relative improvement is 0.02% against a 10% requirement. The paired 95% interval for the *loss difference*, $[-0.0242, +0.0242]$, includes zero. Together with the model-class departures in Section 3.1, this limits the downstream claim: large improvements in derivative accuracy do not establish a consistent hedging advantage specific to derivative supervision in the evaluated settings.

A separate, exploratory zero-cost comparison considers the same coefficient shift. The fully supervised model's hedging loss is 0.7% above that of the same rule using exact Heston delta. Of its remaining gap to an analytic minimum-variance hedge (Föllmer and Schweizer, 1991; Hull and White, 2017), 96.7% is associated with the choice of hedging rule. Models trained with gamma labels reconstruct that hedge from their own derivatives to within the previously specified 2% tolerance; the closest without gamma labels falls 3.3% short (Appendix D). These comparisons are exploratory and leave the registered failure unchanged. At 1% cost, even the residual-only baseline beats exact Heston delta, which is therefore not a lower bound on loss for this task.

Accuracy rankings depend on the evaluation region. Restricting delta RMSE to the hedged contract's region reverses the two baselines: at the interpolated anchor, the residual-only model is better on the full grid (0.0537 versus 0.0881), but worse regionally (0.0937 versus 0.0409). The reversal holds in all five seeds at both anchors (Table 12). This exploratory metric is not trajectory-weighted and does not establish a general predictor of hedging performance.

4 Conclusion

Derivative supervision improves accuracy at held-out Heston coefficients. Adding curvature labels also improves the unsupervised variance derivative. These gains do not establish a consistent hedging advantage specific to derivative labels under the tested departures and costs. The confirmatory test fails, reinforcing the need to evaluate decisions alongside fidelity to the assumed model.

The study covers one two-state PDE, a small network, simulated paths and selected departures, with no real-market validation. Derivative-accuracy comparisons use five seeds; ladder hedging comparisons use ten. Labels and reference share a solver family, so solver agreement cannot exclude common bias. Near zero variance, boundary treatment may affect variance sensitivity. No operator-learning baseline or adaptive residual weighting was tested. The fixed hedge need not be optimal at each transaction cost (Whalley and Wilmott, 1997).

Broader validation should test other PDEs, label-noise levels and combinations of derivative labels. A factorial design would help distinguish each label's effect from its interactions with the others. Whether supervising energy curvature also improves force accuracy in molecular models remains untested here. The practical lesson is to measure both derivative error and decision quality, stating where errors are evaluated, how uncertain the reference is, and which costs the decision incurs.

Acknowledgments and Disclosure of Funding

This work received no specific funding. The authors declare no competing interests.

References

Albrecher, H., Mayer, P. A., Schoutens, W., and Tistaert, J. (2007). The little Heston trap. *Wilmott*, 2007(1):83–92.

Andersen, L. (2008). Simple and efficient simulation of the Heston stochastic volatility model. *Journal of Computational Finance*, 11(3):1–42.

Bates, D. S. (1996). Jumps and stochastic volatility: exchange rate processes implicit in Deutsche Mark options. *The Review of Financial Studies*, 9(1):69–107.

Behler, J. and Parrinello, M. (2007). Generalized neural-network representation of high-dimensional potential-energy surfaces. *Physical Review Letters*, 98(14):146401.

Cannon, P., Ward, D., and Schmon, S. M. (2022). Investigating the impact of model misspecification in neural simulation-based inference. arXiv:2209.01845.

Craig, I. J. D. and Sneyd, A. D. (1988). An alternating-direction implicit scheme for parabolic equations with mixed derivatives. *Computers & Mathematics with Applications*, 16(4):341–350.

Cranmer, K., Brehmer, J., and Louppe, G. (2020). The frontier of simulation-based inference. *Proceedings of the National Academy of Sciences*, 117(48):30055–30062.

Czarnecki, W. M., Osindero, S., Jaderberg, M., Świrszcz, G., and Pascanu, R. (2017). Sobolev training for neural networks. In *Advances in Neural Information Processing Systems (NIPS)*, volume 30.

Ercolessi, F. and Adams, J. B. (1994). Interatomic potentials from first-principles calculations: The force-matching method. *Europhysics Letters*, 26(8):583–588.

Fang, F. and Oosterlee, C. W. (2008). A novel pricing method for European options based on Fourier-cosine series expansions. *SIAM Journal on Scientific Computing*, 31(2):826–848.

Föllmer, H. and Schweizer, M. (1991). Hedging of contingent claims under incomplete information. In Davis, M. H. A. and Elliott, R. J., editors, *Applied Stochastic Analysis*, volume 5 of *Stochastics Monographs*, pages 389–414. Gordon and Breach.

Hainaut, D. and Casas, A. (2024). Option pricing in the Heston model with physics inspired neural networks. *Annals of Finance*, 20(3):353–376.

Heston, S. L. (1993). A closed-form solution for options with stochastic volatility with applications to bond and currency options. *The Review of Financial Studies*, 6(2):327–343.

Huge, B. and Savine, A. (2020). Differential machine learning. arXiv:2005.02347.

Hull, J. and White, A. (2017). Optimal delta hedging for options. *Journal of Banking & Finance*, 82:180–190.

Krishnapriyan, A. S., Gholami, A., Zhe, S., Kirby, R. M., and Mahoney, M. W. (2021). Characterizing possible failure modes in physics-informed neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pages 26548–26560.

Lord, R., Koekkoek, R., and van Dijk, D. (2010). A comparison of biased simulation schemes for stochastic volatility models. *Quantitative Finance*, 10(2):177–194.

Malandain, K. A., Kalici, S., and Chakhoyan, H. (2025). DeepSVM: Learning stochastic volatility models with physics-informed deep operator networks. arXiv:2512.07162.

Merton, R. C. (1976). Option pricing when underlying stock returns are discontinuous. *Journal of Financial Economics*, 3(1–2):125–144.

Noguer i Alonso, M. (2026). Derivative-informed operator learning for finance: On-the-fly Greeks, surfaces, hedging, and control. arXiv:2606.05900.

Raissi, M., Perdikaris, P., and Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707.

Rockafellar, R. T. and Uryasev, S. (2000). Optimization of conditional value-at-risk. *Journal of Risk*, 2(3):21–41.

Ross, S., Gordon, G., and Bagnell, D. (2011). A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 627–635.

Sakuma, T. (2026). Differential machine learning for 0DTE options with stochastic volatility and jumps. arXiv:2603.07600.

Son, H., Jang, J. W., Han, W. J., and Hwang, H. J. (2023). Sobolev training for physics-informed neural networks. *Communications in Mathematical Sciences*, 21(6):1679–1705.

Wang, J., Olsson, S., Wehmeyer, C., Pérez, A., Charron, N. E., de Fabritiis, G., Noé, F., and Clementi, C. (2019). Machine learning of coarse-grained molecular dynamics force fields. *ACS Central Science*, 5(5):755–767.

Wang, S., Teng, Y., and Perdikaris, P. (2021). Understanding and mitigating gradient flow pathologies in physics-informed neural networks. *SIAM Journal on Scientific Computing*, 43(5):A3055–A3081.

Whalley, A. E. and Wilmott, P. (1997). An asymptotic analysis of an optimal hedging model for option pricing with transaction costs. *Mathematical Finance*, 7(3):307–324.

Declaration of AI-assisted technologies

The authors used generative AI and AI-assisted tools for manuscript drafting, language refinement, and code drafting and debugging. The authors directed the project, reviewed and verified the manuscript, the code, and all reported results, and take full responsibility for the final manuscript.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and Section 1 separate the two evidence bases the paper rests on — held-out coefficient regimes inside the model class for the field results, and the downstream cells for model-class departure — state that the registered confirmatory endpoint failed, and scope every claim to this testbed and this nested label ordering. Section 4 states the principal limitations.

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Section 4 discusses the single PDE and small network, simulated rather than real-market data, seed counts, shared solver bias, boundary treatment, missing baselines, and the fixed hedging rule. Section 3.3 explains why the nested design cannot isolate all label effects and interactions. Section 3.5 discloses the unmatched control, and Appendix B reports reference uncertainty.

   Guidelines:

   - The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [N/A]

   Justification: The paper introduces no theorem or new formal proof. The governing equation and the variance-minimising control ratio used in the experiments are stated with their assumptions and cited in Section 2 and Appendix D.

   Guidelines:

   - The answer [N/A] means that the paper does not include theoretical results.
   - All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
   - All assumptions should be clearly stated or referenced in the statement of any theorems.
   - The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
   - Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.

- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

   Answer: [Yes]

   Justification: Section 2 and Appendices A–D disclose the frozen acceptance block, coefficient sampling and anchor excision, the reference construction with its per-derivative uncertainty, architecture, objective, optimizer and schedule, seeds, common-random-number simulation, and the paired-bootstrap analysis needed to reproduce every reported comparison.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
   - If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
   - Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
   - While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
     (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
     (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
     (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
     (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

   Answer: [No]

   Justification: No public code or data archive is linked in this version. The experimental configuration and evaluation procedures are described in the paper and appendices.

   Guidelines:

   - The answer [N/A] means that paper does not include experiments requiring code.
   - Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.

- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: [Yes]

   Justification: Section 2 gives the design and the validation-only selection rule; Tables 6 and 7 give the architecture, optimizer and schedule, batch and collocation sizes, label artifact, seeds, and the per-arm label flags, together with the residual weight and price-point budget each arm actually carried, including the one control that departs from the shared configuration.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
   - The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: [Yes]

   Justification: Section 3 states that path-bootstrap intervals are conditional on the fitted models and that exploratory comparisons are unadjusted for multiplicity. Section 3.6 reports the confirmatory interval; Appendix D gives the paired-bootstrap procedure and identifies point-estimate-only comparisons. Appendix B reports reference uncertainty and unresolved orderings.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
   - The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
   - The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
   - The assumptions made should be given (e.g., Normally distributed errors).
   - It should be clear whether the error bar is the standard deviation or the standard error of the mean.

- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [No]

   Justification: Table 7 reports per-arm wall clock, parameter count, derivative-evaluation count and peak process memory, and Table 6 names the worker type as single-process CPU. The processor model and a project-wide compute total, including preliminary and discarded runs, were not recorded, so the paper does not meet this item's full accounting standard.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
   - The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

   Answer: [Yes]

   Justification: This numerical study uses synthetic data, no human participants or private data, and reports failed registered endpoints and adverse diagnostics rather than suppressing them. AI-assisted drafting is disclosed before the appendices, and the authors retain responsibility for verification.

   Guidelines:

   - The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
   - If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
   - The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [N/A]

    Justification: This is a synthetic methodological evaluation of how derivative labels should be spent when training a surrogate, not a deployed system or a study involving people, so no direct societal impact is claimed. Its results should not be treated as validation of any decision system built on the testbed used here, a scope restriction stated in Sections 1 and 4.

    Guidelines:

    - The answer [N/A] means that there is no societal impact of the work performed.

- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

    Answer: [N/A]

    Justification: The work introduces no pretrained model, scraped dataset, personal data, or other artifact with a high risk of misuse; it consists of small numerical surrogates and synthetically generated trajectories.

    Guidelines:

    - The answer [N/A] means that the paper poses no such risks.
    - Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
    - Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
    - We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: [N/A]

    Justification: The experiments use no external dataset, pretrained model, or third-party research artifact. All labels and trajectories are generated from the equations and numerical methods specified and cited in the paper.

    Guidelines:

    - The answer [N/A] means that the paper does not use existing assets.
    - The authors should cite the original paper that produced the code package or dataset.
    - The authors should state which version of the asset is used and, if possible, include a URL.

- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [N/A]

    Justification: No public asset release is linked in this version. The paper describes the generated labels and trajectories, training configurations, and evaluation procedures.

    Guidelines:

    - The answer [N/A] means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [N/A]

    Justification: The study involves neither crowdsourcing nor human subjects.

    Guidelines:

    - The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
    - Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
    - According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [N/A]

    Justification: The study involves no human subjects, so institutional human-subjects review is not applicable.

    Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

    Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

    Answer: [N/A]

    Justification: LLMs are not an important, original, or non-standard component of the core scientific method. The paper nevertheless voluntarily discloses their use for manuscript and code drafting and debugging in the declaration preceding the appendices; the authors reviewed and verified the work.

    Guidelines:

    - The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
    - Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

A Pre-registered acceptance block and amendment record

The decision-bearing portions below are transcribed from `heston_benchmark_v6.yaml`, the contract consumed by the frozen label and training artifacts. Long audit comments are omitted, but key names and values are unchanged. The label artifact is frozen at tag `v6-labels-20260812`.

```
acceptance_thresholds:
  confirmatory_cell_pass: "misspec delta CVaR95 improvement >=10% relative AND paired-bootstrap 95% CI excludes 0
      (combined,1% TC,baseline,rung3 vs standard_pinn)"
  confirmatory_cell_rel_min: 0.10
  headline_scale_free: fraction_of_baseline_to_oracle_gap_closed
  order_attribution: "rung2 beats rung1 at the cell, CI excludes 0; failure = honest null"
  dose_response:
    criteria: "Spearman rho > 0 (shape confirmed via isotonic fit) AND one-sided seed-bootstrap tail probability P(rho<=0)
      < bootstrap_tail_prob_max"
    bootstrap_tail_prob_max: 0.05
    interpretation: "monotonicity not demonstrated is not the claim that the response is flat"
  ood_gamma_rmse_reduction_min: 0.15
  ood_vega_rmse_reduction_min: 0.15
  price_parity_within: 0.10
  in_model_hedging: NOT_PASS_FAIL
  sakuma_null_rel_tol: 0.02
  seeds: {default: 5, confirmatory_cell: 10}
  tail_claim_requires: paired_bootstrap_over_CRN_paths_with_seed_variance_separated

lambda_selection:
  lambda_delta: {value: 1.0, fixed: true, swept: false}
  lambda_pde:
    source_arm: standard_pinn
    candidates: [0.0, 0.01, 0.1, 1.0]
  lambda_gamma:
    source_arm: rung3_delta_gamma_vega
    candidates: [0.3, 1.0, 3.0]
  lambda_vega:
    source_arm: rung3_delta_gamma_vega
    candidates: [0.3, 1.0, 3.0]
  tune_on: validation_only
  robustness_row:
    description: "also report the confirmatory contrast at the rung3-sourced lambda_pde"
    status: robustness_result_not_a_second_confirmatory_test
```

The selected weights were $(\lambda_{\text{pde}}, \lambda_D, \lambda_G, \lambda_v) = (0.01, 1, 1, 1)$. The registered alternative sourced from rung 3 selected $\lambda_{\text{pde}} = 0$ and was reported as a robustness row.

Amendment record. The contract evolved through four documented, pre-run amendments. All preceded training and model evaluation; amendment 3 used a production-scale *protocol diagnostic* to correct the clipping rationale, but no trained-arm outcome.

Table 2: Contract amendments before the frozen label artifact. "Pre-eval" means before any trained-arm evaluation.

ID	Date	Pre-eval	Reason and effect
A1	2026-07-31	yes	Declared formerly implicit thresholds; recorded the realised 43-step clock; distinguished "monotonicity not demonstrated" from flatness; fixed validation-only weight sourcing.
A2	2026-08-04	yes	Made the 10% confirmatory threshold numeric; separated absent from broken evaluations in the verdict vocabulary; declared the corruption-noise ladder of the pre-training detectability diagnostic, and its region of validity.
A3	2026-08-05	yes	Corrected the measured clipping direction in that same diagnostic, demoted its 0.20 noise rung to diagnostic-only, and fixed the extension contingency. The triggering data were clip measurements, not arm outcomes.
A4	2026-08-05	yes	Declared the 20% parameter-point validation split, 64 state–contract–time draws per parameter point, and full Monte Carlo coverage.

Table 3: Registered endpoints and their outcomes. "Not demonstrated" is the contract's own wording for a dose-response that fails its monotonicity criterion, which is not the claim that the response is flat. The in-class check is declared `NOT_PASS_FAIL`: it flags a registered expectation motivated by the similar pricing and Greek accuracy of the supervised models in Sakuma (2026), not a hedging null established by that paper. Intervals below are in loss-difference units; the dose probability is a one-sided seed-bootstrap tail probability.

Endpoint	Declared	Observed	Verdict
Confirmatory: misspecified J, rung 3 vs baseline	$\geq 10\%$ relative, CI excludes 0	$+0.02\%$, CI $[-0.0242, +0.0242]$	fail
Order attribution: rung 2 beats rung 1	CI excludes 0	$+0.0574$, CI $[+0.0527, +0.0621]$; -1.2% relative	fail (reversal)
Dose-response monotonicity	$\rho_S > 0$ and $P(\rho_S \leq 0) < 0.05$	$\rho_S = 0.029$, $P = 0.074$	not demonstrated
Held-out $\partial_{SS}u$ and $\partial_v u$ RMSE at value parity	$\geq 15\%$ each, value parity within 10%	reductions 0.864–0.930; value improved 82–91%	fail (on parity)
In-class consistency check	not pass/fail	$+27.2\%$ gap at zero cost	flag
Mechanism rule	zero-cost gap or cost/-turnover pattern	zero-cost gap; later controls do not isolate a Gamma-specific cause	channel (i), qualified
Departure sweep: locate a favourable regime	—	none at the 1% tier	no decisive regime

B Reference construction and per-derivative solver spread

Section 2 reports agreement across available numerical legs. For derivatives, the measured spread is a diagnostic scale for arm comparisons; it is not a bound on reference error or a guarantee of an ordering.

Construction, and which legs run where. Leg 1 analytically differentiates the trap-free Heston P_1/P_2 Fourier integrands (Heston, 1993; Albrecher et al., 2007) before Gauss–Legendre quadrature. Leg 2 differentiates a Fang–Oosterlee cosine-expansion pricer (Fang and Oosterlee, 2008) with fourth-order central stencils and a three-step Richardson sweep. It separately implements the same ϕ_2 formula as leg 1, so their agreement cannot exclude a shared formula error. Leg 3 is Monte Carlo, pathwise for the first derivative and mixed pathwise/likelihood-ratio for curvature, with antithetic pairs and a control variate; it returns standard errors and has no estimator for the mixed derivative. Leg 4 is a Craig–Sneyd ADI solve (Craig and Sneyd, 1988), routed only where the variance

process is stressed. The interpolated anchor (Feller ratio 1.04) draws all four legs and a four-fold Monte-Carlo path multiplier, $800{,}000$ paths; the extrapolated anchor (Feller ratio 1.78) draws three, $200{,}000$ paths, and has no ADI leg. The legs have distinct numerical procedures but do not provide four fully independent certifications.

Per grid point the reference value is the median across legs and the reported uncertainty is the largest pairwise absolute disagreement, folded with the Monte-Carlo standard error; a point whose largest disagreement exceeds 10^{-3} times a wing-floored scale, plus $3\sigma_{\mathrm{MC}}$ for any pair involving the stochastic leg, is masked. The rule was declared before training.

Deterministic spread and stochastic uncertainty. Table 4 reports RMS disagreement among deterministic legs alongside the broader envelope. At the extrapolated anchor the two characteristic-function legs agree to 1.6×10^{-8} on $\partial_S u$ and 3.3×10^{-8} on $\partial_{SS} u$ in RMS. This agreement cannot rule out their shared bias. The curvature envelope is 2.65 times larger there despite the interpolated anchor's lower Feller ratio. The stored Monte-Carlo standard errors have ratio 2.63; a nominal $n^{-1/2}$ adjustment for the fourfold path-budget difference reduces this to 1.32. This is a scaling diagnostic, not an equal-budget rerun or a bias/convergence certificate. The earlier pointwise MC-dominance percentages are omitted because an audit did not consistently reproduce them under explicit dominance definitions.

Table 4: Per-derivative reference uncertainty at the two evaluation anchors, in the units of Table 1 (curvature $\times 10^{-3}$), as root mean square over the retained points each arm is scored on. "Envelope" is the quantity the paper reports: largest pairwise leg disagreement folded with the Monte-Carlo standard error. "Deterministic" is the largest pairwise disagreement among the non-stochastic legs alone, recomputed on the frozen grids. The ratio is envelope divided by deterministic spread; neither quantity bounds common solver bias.

Regime	Derivative	envelope	deterministic	ratio
interpolated	$\partial u/\partial S$	0.000276	6.96×10^{-6}	40
(4 legs, 8×10^5 paths)	$\partial^2 u/\partial S^2$	0.768	0.00156	491
	$\partial u/\partial v$	0.1083	0.000656	165
extrapolated	$\partial u/\partial S$	0.000397	1.61×10^{-8}	2.5×10^4
(3 legs, 2×10^5 paths)	$\partial^2 u/\partial S^2$	2.035	3.31×10^{-5}	6.1×10^4
	$\partial u/\partial v$	0.1270	5.07×10^{-6}	2.5×10^4

Limits of the resolution heuristic. The reference is a median across available legs. Its agreement with characteristic-function values does not prove truth or independence. Even a valid RMS reference-error bound ϵ would allow a difference of two RMSEs to move by up to 2ϵ; measured solver spread need not be such a bound. Exceeding one envelope therefore supplies a sensitivity diagnostic, not a one-sided guarantee. Falling below it is likewise not evidence of equality.

Which contrasts in Table 1 survive. Table 5 divides each arm-to-arm RMSE difference by the envelope for that derivative and anchor. The convention follows the study's registered `improvement_to_oracle_noise` statistic; it is a resolution heuristic, not a hypothesis test, and it ignores seed dispersion, which is reported separately in Table 1's source.

Table 5: Arm-difference divided by the reference envelope of Table 4. Positive means the second arm is more accurate. A magnitude below 1 falls below the declared resolution heuristic; this is not an inferential significance threshold.

	interpolated			extrapolated		
Contrast	$\partial_S u$	$\partial_{SS} u$	$\partial_v u$	$\partial_S u$	$\partial_{SS} u$	$\partial_v u$
baseline → rung 1	156.2	6.9	67.1	86.8	2.2	45.0
baseline → rung 3	184.7	9.1	85.0	97.3	2.6	54.5
rung 1 → rung 2	26.2	2.0	9.9	11.4	*0.5*	3.5
rung 2 → rung 3	2.2	*0.2*	8.0	*−0.97*	*−0.1*	6.0

Three readings follow, and Sections 3.1–3.4 are written to them.

The headline survives everywhere. Every baseline-to-supervised contrast exceeds the envelope at both anchors, by factors from 2.2 to 185. The 82 to 93% reductions exceed this diagnostic scale, conditional on the reference.

The curvature rung is resolved at one anchor only. The rung-1-to-rung-2 curvature gap is 2.0 envelopes at the interpolated hole and 0.5 at the extrapolated one. Section 3.2 claims the increment at the interpolated anchor and declines to read its extrapolated magnitude finely. The transfer result of Section 3.3 lives in the $\partial_v u$ column and is resolved at both anchors, at 9.9 and 3.5 envelopes.

The two extrapolated reversals are not resolved, and the appendix says so plainly. The flagged risk was that Section 3.4's non-monotonicity rests on $0.641 \to 0.823$ on $\partial^2 u/\partial S^2$ and $0.003399 \to 0.003784$ on $\partial u/\partial S$. Those gaps are 0.09 and 0.97 envelopes. Both sit inside the reference uncertainty, so this benchmark cannot call either a degradation, and Section 3.4 does not. What survives there is the resolved gain on the derivative rung 3 adds, at 8.0 and 6.0 envelopes.

A sharper form of the same limit: at the extrapolated anchor the curvature RMSE of rungs 1, 2 and 3 is 0.77, 0.31 and 0.40 times the envelope. All three errors fall below the envelope, so fine orderings remain conditional on the reference rather than independently certified. Only the gap from the unsupervised arms, which are 3.0 and 3.5 envelopes away, is.

Masking: rate, location, and one correction. The union mask over all five reference quantities removes 17.11% of the interpolated grid (3,704 of 21,648 points) and 9.23% of the extrapolated grid (1,999 of 21,648). Per quantity the interpolated rates are 0.008 (value), 0.038, 0.073, 0.080 and 0.116 for the first derivative, curvature, second state and mixed derivative; the extrapolated rates are 0.016, 0.014, 0.055, 0.044 and 0.052. Table 1 scores every arm on the same union mask.

The mask is a relative-tolerance artefact and sits where that tolerance is hardest to meet. It concentrates overwhelmingly at the shortest maturity, 94.3% of the $\tau = 0.04$ slice at the interpolated anchor and 78.7% at the extrapolated one, against 0.077 to 0.102 beyond $\tau = 0.45$ at the interpolated anchor and 0.000 to 0.020 beyond $\tau = 0.29$ at the extrapolated one; and it peaks just out of the money rather than deep out of it, 47.5% in the $S/K \in [0.75, 0.9)$ band at the interpolated anchor against 18.1% below 0.75 and 6.8% above 1.3. In curvature the pattern is U-shaped, heaviest in the lowest decile, where a relative tolerance on $|\partial^2 u/\partial S^2| \leq 3 \times 10^{-12}$ demands absolute agreement no solver delivers, and in the highest.

The extrapolated anchor is less masked, 9.23% versus 17.11%. The fourth leg at the interpolated anchor introduces additional pairwise checks. These masks are selective; shared support makes arm comparisons consistent on retained points but does not establish performance over excluded states.

C Training configuration and excision protocol

The main ladder and baseline arms share architecture, optimiser, data and step budget; their loss terms differ as listed below. The extra-value-data sweep is an exception. Table 6 states the shared budget once and Table 7 the per-arm difference, so that the identical-configuration claim of Section 2 is checkable rather than asserted.

Excision. Training coefficients are a seeded Latin hypercube on $\kappa \in [1, 4]$, $\theta \in [0.02, 0.12]$, $\xi \in [0.20, 0.60]$, $\rho \in [-0.80, -0.20]$, $v_0 \in [0.01, 0.12]$, rejecting any draw with $2\kappa\theta/\xi^2 < 0.40$ (an oracle-trust cutoff, not the Feller condition). Around each of the five named regimes a ball is excised: in coordinates range-normalised to the unit cube, every draw within Euclidean distance $0.10\sqrt{5} = 0.2236$ of an anchor, that is 10% of the cube diagonal, is rejected. Realised: no sampled point falls inside any ball, the nearest lying at 1.03 to 1.64 times the radius, minimum normalised anchor distance 0.2310 to 0.3670; minimum realised Feller ratio 0.4061.

The validation split is 20% of *coefficient points*, not of rows, so a held-out regime cannot leak into validation through a shared coefficient draw. Both evaluated regimes are therefore absent from training and from every selection score. The interpolated anchor $(\kappa, \theta, \xi, \rho, v_0) = (1.5, 0.04, 0.34, -0.70, 0.02)$ sits inside the sampled box in every coordinate and is a hole punched in the support. The extrapolated anchor $(2.0, 0.04, 0.30, -0.90, 0.04)$ has $\rho = -0.90$ against a sampled range whose lower bound is -0.80 and whose minimum realised draw is -0.7989: the ball

Table 6: Shared configuration. Every entry is identical across the ladder and baseline arms; the extra-value-data control of Table 7 departs from it in the weight, row-budget and selection settings described below.

Component	Setting
Inputs	$(S, K, \tau, \kappa, \theta, \xi, \rho, v_0)$, affinely mapped to $[-1, 1]$ by fixed contract ranges inside the forward pass with no clamp, so evaluated regimes outside those ranges extrapolate rather than saturate
Network	4×64 `tanh` multilayer perceptron, 13,121 parameters, `float32`; no hard-constrained ansatz, terminal treatment identical across arms
Optimiser	Adam, learning rate 3×10^{-3} cosine-annealed to 3×10^{-5}, gradient-norm clip 1.0
Budget	20,000 steps, minibatch 1,024 label rows, 4,096 residual collocation points redrawn each epoch (834 epochs); fixed-step mode (called matched-epoch in code), early stopping disabled
Labels	23,623 training and 5,936 validation rows from the frozen artifact `v6-labels-20260812`: 512 coefficient draws crossed with 64 shared (S, K, τ) triples, mask rate 0.098
Loss normalisation	supervised terms divided by second moments computed once over the retained training split; the Gamma scale uses true consensus `gamma_ref`, shared across arms. Without it the curvature term sits about six orders of magnitude below the value term at equal relative error and is numerically inert
Weights	$(\lambda_{\mathrm{pde}}, \lambda_1, \lambda_2, \lambda_v) = (10^{-2}, 1, 1, 1)$, selected on validation only, sourced per Appendix A
Seeds	5 per arm; 10 on the ladder arms and the residual-only baseline
Hardware	CPU, single process; per-arm wall clock and peak memory in Table 7. Processor model and a project-wide compute total were not recorded

Table 7: Arms: label flags and measured cost. *res.* is the physics residual; ∂_S, ∂_{SS}, ∂_v are the three label terms. Wall clock is the per-seed mean; peak memory is the seed-42 measurement. Derivative evaluations count first- and second-order AD calls, rather than only double-backward passes. $*$ The extra-value-data control uses residual weight 1 instead of 10^{-2} and 12,288 retained rows instead of 23,623. Its checkpoint and run log have conflicting configurations; its cost figures are withheld.

Arm	res.	∂_S	∂_{SS}	∂_v	seeds	wall (s)	deriv. evals	peak (MB)
value-only feedforward	–	–	–	–	5	17	0	383
residual-only baseline	✓	–	–	–	10	377	60,120	978
rung 1, $+\partial_S u$	✓	✓	–	–	10	414	80,160	941
rung 2, $+\partial_{SS} u$	✓	✓	✓	–	10	471	100,200	1009
rung 3, $+\partial_v u$	✓	✓	✓	✓	10	451	100,200	1043
supervised, residual-free	–	✓	✓	✓	5	64	40,080	695
curvature L^2 penalty	✓	–	–	–	5	465	100,200	1267
shuffled labels	✓	✓	✓	✓	5	433	100,200	1095
extra value data*	✓	–	–	–	5	—	—	—

intersects the cube only where it overlaps, and the anchor itself lies outside the sampled support entirely, so that regime is extrapolation and not interpolation into a hole.

The contract grid is $41 \times 33 \times 16$ over $S \in [50, 150]$, $K \in [60, 140]$, $\tau \in [0.04, 1]$ at $r = 0.02$, $q = 0$. Training labels use 64 shared (S, K, τ) triples, giving sparse local coverage near the hedged contract. Absence of a nearby sampled triple does not by itself establish extrapolation outside the training support in these coordinates.

Departure variants. The Bates generating process adds a compound-Poisson component (Bates, 1996) to the baseline anchor $(\kappa, \theta, \xi, \rho, v_0) = (2.0, 0.04, 0.30, -0.50, 0.04)$, over $\lambda_j \in \{0, 0.1, 0.25, 0.5\}$ and $\sigma_j \in \{0.05, 0.10, 0.15\}$ at $\mu_j = -0.10$. The Merton variant (Merton, 1976) replaces the two-state diffusion with a one-state jump-diffusion at $\sigma = 0.20$, $\lambda_j = 0.25$, $\mu_j = -0.10$, $\sigma_j = 0.10$, so the surrogate stays conditioned on a second state the world does not have. Every arm is handed the baseline anchor's coefficients at decision time, never the true ones. The $\lambda_j = 0$ cell is the contract's in-class consistency check and reproduces the Heston cell exactly, 1.7371 against 1.7371. It is the least-departed cell in Table 8 and carries the largest advantage in

it, $+27.7\%$ at zero cost. A second registered check compares the same corner with an expectation of near-indifference motivated by the similar pricing and Greek accuracy of the three supervised models in Sakuma (2026). It *flags*: the gap exceeds that registered expectation. Their accuracy comparison does not itself predict a hedging null in our setting. We retain the flag as a declared consistency check, not a pass/fail endpoint.

The registered severity sweep, read out. Table 8 gives the sweep Section 3.1 summarises. At zero action cost the advantage of the fully supervised arm over the residual-only baseline falls monotonically in *each* departure coordinate, all six one-dimensional chains decreasing, from $+27.7\%$ at the in-class cell to -4.2% at the most departed one, changing sign between $\lambda_j = 0.5,\ \sigma_j = 0.05$ and $\lambda_j = 0.5,\ \sigma_j = 0.10$. Every pooled interval excludes zero except the sign-crossing cell itself. At the registered 1% cost the whole table is harm and no cell is favourable, which is what produced the registered `no_decisive_regime` verdict: that verdict reads only the 1% tier. The ordering is ordinal. We know which cells are more departed, and the advantage decays in that order; we did not compute a distance from the assumed model class, so the axis has no calibrated scale and the table is not a severity-versus-advantage *curve* in the metric sense.

Table 8: Registered departure sweep, delta-only J (lower is better), five seeds, shared trajectories. "Advantage" is the relative improvement of rung 3 over the residual-only baseline; positive favours supervision. † marks the one cell whose pooled paired-bootstrap interval on the difference includes zero. The $\lambda_j = 0$ row is the in-class consistency check, not a departure.

		zero cost		1% cost	
λ_j	σ_j	rung 3	advantage	rung 3	advantage
0 (in class)	0.05	1.737	$+27.7\%$	4.818	-6.9%
0.10	0.05	2.096	$+20.5\%$	5.082	-7.8%
0.10	0.10	2.358	$+17.9\%$	5.330	-7.5%
0.10	0.15	2.699	$+15.6\%$	5.669	-7.0%
0.25	0.05	2.723	$+10.6\%$	5.533	-9.1%
0.25	0.10	3.455	$+7.0\%$	6.207	-8.1%
0.25	0.15	4.384	$+4.9\%$	7.113	-6.5%
0.50	0.05	3.582	$+0.8\%^{\dagger}$	6.205	-11.1%
0.50	0.10	5.057	-3.0%	7.532	-9.6%
0.50	0.15	6.870	-4.2%	9.280	-7.8%
state removed		2.925	$+9.8\%$	5.845	-3.7%

D Supporting exhibits

Value accuracy does not imply derivative accuracy. Table 9 is the exhibit behind Section 1's opening clause. Both unsupervised arms match the reference solution to a few percent while their curvature is wrong by tens of percent, and the ordering on the value is not the ordering on the curvature: the residual-only arm has the better value fit and the worse curvature.

Table 9: Relative RMSE against the reference, at the two held-out coefficient regimes, five seeds. Controlling the value in L^2 leaves curvature error one to two orders of magnitude larger in relative terms.

	interpolated		extrapolated	
Arm	u	$\partial^2 u/\partial S^2$	u	$\partial^2 u/\partial S^2$
residual-only baseline	0.0155	0.4096	0.0128	0.3505
value-only feedforward	0.0328	0.2664	0.0318	0.2206
rung 3 (fully supervised)	0.0014	0.0446	0.0022	0.0498

The information-free controls. Table 10 gives the three controls of Section 3.5 with seeds and dispersion, together with the first-derivative column, which is what separates "no signal" from "objective conflict": the curvature penalty is worse than the baseline on the derivative it never touches

as well as on the one it shapes. The two configuration-matched controls are the ones the argument rests on; the third is discussed below.

Scope of the extra-value-data sweep. The saturation driver uses $\lambda_{\mathrm{pde}} = 1$ rather than the selected 10^{-2}. Its nominal base $N = 4{,}096$ is also smaller than the main arms' actual 23,623 retained rows: even $5N = 20{,}480$ remains below the main budget. The selected $3N$ checkpoint uses 12,288 rows. Thus the level comparison in Table 10 is not matched in weight or data.

Within the sweep, validation curvature relative RMSE changes from 0.647 ± 0.095 to 0.596 ± 0.015 across $1N$–$5N$. The declared 2% rule selects $3N$ (0.605 ± 0.017); its doubled-width control gives 0.709 ± 0.061. This describes the tested protocol and does not rule out optimization limitations. The sweep uses its own validation-selection protocol; its scores are distinct from held-out RMSE. A substitution claim would require matched actual row budgets, residual weights and selection rules. The penalty and shuffled-label controls are configuration-matched, but their failures do not prove a general information bound.

Table 10: Control arms, held-out relative RMSE, mean ± seed standard deviation over five seeds. λ_{pde} and N are the residual weight and actual retained price-row count. Curvature relative RMSE uses the declared central moneyness support; absolute RMSE in Table 1 uses the full retained grid. The extra-data row is not a matched comparison.

			$\partial^2 u/\partial S^2$ rel. RMSE		$\partial u/\partial S$
Arm	λ_{pde}	N	interpolated	extrapolated	interpolated
residual-only baseline	0.01	23,623	0.4096 ± 0.0106	0.3505 ± 0.0158	0.0749
curvature L^2 penalty, no labels	0.01	23,623	0.9822 ± 0.0068	0.9776 ± 0.0079	0.5389
shuffled labels	0.01	23,623	0.4075 ± 0.0343	0.3666 ± 0.0300	0.1360
extra value data (plateau)	1.00	12,288	0.6986 ± 0.0118	0.6147 ± 0.0116	0.2108
rung 3 supervision	0.01	23,623	$\mathbf{0.0446 \pm 0.0028}$	$\mathbf{0.0498 \pm 0.0018}$	**0.0038**

Reconstructing the variance-minimising controller. The combination $\partial_S u + (\rho\xi/S)\,\partial_v u$ is the local-risk-minimising ratio for this state (Föllmer and Schweizer, 1991). Any arm carrying both first derivatives can assemble it from its own outputs with no further reference, so Table 11 is the point at which the field result and the decision result meet. It is exploratory and post-registered, moves no registered verdict, and reuses the frozen checkpoints and the same shared trajectories every other cell faces; ρ and ξ in the correction are the assumed-model values, so no true-process parameter enters. Point estimates only: these are five-seed means against a pre-registered 2% floor, and we claim no interval for them.

Both arms carrying curvature labels land inside the floor, as does the residual-free supervised arm; the two arms without curvature labels miss it by 3.3 and 3.4 percent, and the residual-damaged baseline by 42.8. The reference the registered metric scores against, the exact first derivative alone, is itself 24.7% worse than this controller, which is the sense in which the registered endpoint could not have rewarded accuracy at the accuracy supervision reaches.

Table 11: Delta-only J (lower is better) when the controller is built from each arm's own two first derivatives, on the confirmatory cell's combined perturbation at zero action cost, five seeds, shared trajectories. The gap is relative to the analytic construction; negative beats it. The pre-registered negligibility floor is 2%.

Fields driving the controller	J	gap	within 2% floor
analytic construction	1.6696	—	reference
rung 2 ($+\partial_{SS}u$)	**1.6563**	-0.79%	yes
rung 3 ($+\partial_v u$)	1.6700	$+0.03\%$	yes
supervised, residual-free	1.6793	$+0.58\%$	yes
rung 1 ($+\partial_S u$)	1.7240	$+3.26\%$	no
value-only feedforward	1.7261	$+3.38\%$	no
residual-only baseline	2.3843	$+42.81\%$	no
exact first derivative alone	2.0816	$+24.68\%$	no

Delta accuracy on the full grid and the hedging region. The registered accuracy metric uses the full grid, $S \in [50, 150]$, $K \in [60, 140]$, $\tau \in [0.04, 1]$. The hedge trades one contract, with $K = 100$ and remaining expiry $\tau \in [0.08, 0.25]$. Table 12 restricts the same metric to the corresponding strike, expiry interval and the contract's predefined central range $S/K \in [0.75, 1.30]$. It is a restricted-grid comparison; it does not weight states by how frequently trading trajectories visit them.

As summarised in Section 3.6, delta accuracy reverses the two models without derivative labels: the residual-only model is better on the full grid and worse in the restricted region, in all five seeds at both regimes. Gamma also reverses in the mean, but its full-grid ordering holds in only three of five and four of five seeds. We therefore highlight the delta result in the main text. Across training labels in the same expiry and S/K ranges, pooling strikes, the mask rate is 1.30 times the rate outside. This indicates thinner reference-label coverage in that part of the domain.

Table 12: Delta RMSE for the two models without derivative labels, on the full parameter grid and the restricted hedging region, five seeds. Lower is better. The rankings reverse; the regional metric is not trajectory-weighted.

	interpolated		extrapolated	
Arm	full grid	region	full grid	region
residual-only baseline	**0.0537**	0.0937	**0.0424**	0.0689
value-only feedforward	0.0881	**0.0409**	0.0815	**0.0227**